

Leveraging Image Retrieval for Unsupervised Detection of Recurring Content in TV Shows' Video Files

Master Thesis

Niels ten Boom
s4767314

Abstract

This thesis presents research on image retrieval-based approaches for recurring content detection in Video On Demand (VOD) TV-shows. The method utilizes feature vectors constructed from video frames and an image retrieval framework to efficiently compare similarities across TV-show episodes. With such methods the best precision score achieved was 0.884 and the best recall score achieved was 0.892 for recurring content, an altered experimental setup for bumper detection achieved a best precision score of 0.611 and a best recall score of 0.979. With these results we present novel research for unsupervised detection of recurring content in VOD content.

Contents

1	Introduction	4
1.1	Recurring Content Classes	5
1.1.1	Recaps	5
1.1.2	Opening credits	5
1.1.3	Bumpers	5
1.1.4	Previews	6
1.1.5	Closing Credits	6
1.2	Definitions and Techniques	6
1.2.1	Shot Change Detection	6
2	Related Work	7
3	Image Retrieval	8
3.1	Feature vectors	9
3.2	Color Histograms	11
3.3	Color Texture Moments	11
3.4	CNN Features	12
4	Methodology	13
4.1	Data	13
4.2	Feature Vector Construction	14
4.2.1	Color Histograms	14
4.2.2	Color Texture Moments	14
4.2.3	CNN Features	15
4.3	Matching and Detection	15
4.3.1	Recurring Content	17
4.3.2	Bumpers	17
5	Results	17
5.1	Recurring Content	18
5.2	Bumpers	20
6	Discussion	20
7	Conclusion	21
8	References	22

1 Introduction

Viewing rates for TV are dropping gradually while the number of users of streaming services such as Netflix and Videoland are growing rapidly year over year. This shift from TV to Video on Demand introduces new problems for broadcasters, as they now have to support a hybrid form of traditional broadcasting and Video On Demand (VOD). In this thesis we consider a method that could help to improve the user experience for VOD.

Video content is optimized for just one form, and currently that is TV. This means that the video content may contain recurring content consisting of: recaps, opening credits, bumpers (to ease a viewer into commercials), closing credits and previews. A viewer watching this content back-to-back may find it preferable to be able to skip this recurring content to improve their viewing experience. Providing this functionality to users, requires metadata on where the skippable content occurs in the videos. For a large percentage of content, such metadata has not been retained. Videoland currently hosts over 1000 different titles consisting of multiple seasons and episodes and this selection is always changing. An automated solution that can detect this content in videos would be very useful to transform the video content into a format suited for VOD. In this thesis we will explore methods for such an automated solution.

There is little margin for error with this solution. The accuracy of the end result should be sufficiently high enough for it to be trusted to label all of the content automatically. If that is not the case then it needs double checking by a human. Because if a part of the video gets mislabeled, a user may skip over actual content. It is critical to RTL that this does not happen.

This thesis presents research on unsupervised detection of recaps, opening credits, bumpers, closing credits and previews given the video files from a TV-show. We are going to attempt to tackle this problem by using image retrieval techniques. The motivation behind this is that the overlapping characteristic of the segments is that they reoccur. To be able to detect recurring content we are going to compare similarities of frames across content, this should be done efficiently and accurately, both attributes are important in image retrieval research. Our main research question then is:

How well do image retrieval methods perform in accurately detecting recurring content across a TV-show?

With the subquestions:

How can image retrieval be used to detect recurring content?

How accurate is each image retrieval approach?

How can the best method be utilized in practice?

This thesis aims to answer aforementioned research questions. Our hypothesis is that such a method works very well. The rest of this section will be used to expand on the different type of video classes which this research is focused on. Section 2 explains all

of the related work. In Section 3 image retrieval and the used methods are expanded upon. Section 4 describes the data and methodology and then the produced results are presented in Section 5. Lastly, the results are discussed and a conclusion is drawn in Section 6.

1.1 Recurring Content Classes

In this section we will give a background on the different types of content to be detected. What all the content classes have in common is that they (partially) reappear in previous or future episodes or as is the case with bumpers: that the content reappears within the episode itself. We will use this attribute for the detection. Example frame grabs from each of the content classes can be viewed in Figure 1.



Figure 1: Frame grabs from each recurrent content class.

1.1.1 Recaps

A television show typically has recaps before the opening credits. In the recaps content of previous episodes is repeated to refresh the viewer's memory. This is useful for linear TV when an episode is aired every week. Someone watching episodes back-to-back ideally wants to skip the recaps together with the opening credits because they have seen the content recently. Not all material preceding the opening credits is a recap though, it can be original content.

1.1.2 Opening credits

The opening credits of a TV-show are generally the same for all of the episodes in the same season. It typically opens the show with a theme song, presenting the most important actors at the start. If the sequence is known beforehand then it would be a more simple problem to solve, by matching this sequence with all of the videos to locate the opening credits. But there is also no prior knowledge on how the opening credits of a show look like and they often change every season.

1.1.3 Bumpers

The bumper is a part of the video that eases a viewer into the upcoming commercial. Most of the times it has a voice over and text on screen saying: 'Next' (or the dutch translation of 'next') and some footage of what is to be expected after the commercial break is shown.

1.1.4 Previews

A preview is a segment at the end of an episode where an advance showing of fragments of the next episode(s) are played. Previews can typically be found in content that was created specially for broadcast TV. It gives the viewer a taste of what to expect in the next episode that will air a week later. These previews may contain overlaid credits of people who helped in the production of the show.

1.1.5 Closing Credits

The closing credits at the end of a video contain a lot of scrolling text most of the time and there is little variation between the frames. In these texts everybody related to the production of the video is mentioned. In general the frames all have a black background with white text. But backgrounds can also vary, therefore a solution that simply looks for black and white would not suffice.

1.2 Definitions and Techniques

This section will be used to elaborate on some topics mentioned in this thesis. We define a season with multiple episodes of a TV-show as T .

$$T = \{E_1, E_2, \dots, E_x\}$$

An episode within a season is defined as a set of frames.

$$E = \{f_1, f_2, \dots, f_y\}$$

The videos contain shot boundaries, we refer the set of locations of shot boundaries as B .

$$B = \{b_1, b_2, \dots, b_z\}$$

1.2.1 Shot Change Detection

Shot detection is a technique that can be used on videos to determine shot boundaries, see Figure 2 for an example of such a shot change.




Figure 2: Example of a shot change within 4 frames of a video.

Shot detection is going to be used and thus will we expand on it. A lot of research has been done related to shot detection [1] and many techniques have been proposed.

We use a method that looks at the shift in mean color in HSV space [2]. HSV (Hue, Saturation, Value) color space is a more intuitive color mixing model compared

to RGB (Red, Green, Blue) color space. Because one can change a value in either of the three values in HSV and expect what the new color is going to look like. This is almost impossible for RGB because for this same color change to happen, you need to change all three red, green and blue values to result in the same color space.

The shot boundaries are classified by calculating the difference in hue, saturation and value for each frame and then from this the mean is calculated. If the mean is higher than a threshold H , then a shot boundary is marked, refer to algorithm 1 for a pseudo-code representation.

Algorithm 1: Shot boundary detection

```

shotBoundaries = List();
previousMeanHSV = getMeanHSV(Episode[0]);
for frame in Episode do
    meanHSV = getMeanHSV(frame);
    if  $abs(meanHSV - previousMeanHSV) > H$  then
        | shotBoundaries.add(frame);
    end
    previousMeanHSV = meanHSV;
end

```

2 Related Work

No prior work exists that explores a solution for the problem previously described (unsupervised detection of these specific recurring segments). Related work exists that focuses on commercial or repeated sequence detection in large broadcast streams. These detection methods can be roughly divided into three groups: fingerprint, feature-based and unsupervised methods.

Fingerprint methods set up a database of fingerprints of known commercials or repeated sequences to detect these in a broadcast stream. Lienhart et al. [3] propose a method based on features to roughly localize advertisements in a stream and then construct a fingerprint database based on color coherence vectors. Gauch et al. [4] propose a method based on constructing feature vectors from color moments in a stream. Covell et al. [5] implement a fingerprinting method based on audio with visual verification after a proposed match. The disadvantage of these fingerprinting methods is that it is not unsupervised and the setting up and maintenance of such a fingerprint database requires much work.

The feature-based methods detect commercials based on extracted video features. Wang et al. [6] fuse the results of audio scene changes and textual content similarity between shots to segment programs including commercials.

With the unsupervised methods the authors typically do a dimensionality reduction operation and then try to find repeated sequences or commercials with clustering methods. Herley et al. [7] convert the stream with a Discrete Cosine Transform (DCT) to reduce dimensionality and then propose an extensive probability framework to detect repeated sequences. Benezeth et al. [8] use the Electronic Programme Guide (EPG) in addition to the dimensionality reduction to detect program boundaries.

Abduraman et al. [9] propose a system to detect repeated sequences in streams by performing a DCT operation on all of the frames in the stream and then use a micro clustering technique to detect repeated sequences. They were also able to link the trailers to their respective programs that occur at a later point in the stream.

All of the previously mentioned works in this section never cover the specific topic of segmenting the classes mentioned in Section 1.1 or use other properties only available to broadcast streams such as the EPG. The methods also focused on recurrent content detection are exclusively based on detection of moderate to long sequences, it does not cover short recurrent sequences stitched together such as a recap. The methods yielded high accuracies in the range of 90% - 95% precision, also most research is from many years ago and work that focuses on repeated sequence detection in broadcast streams do not vary much in methodology. This work is not focused on broadcast streams and aims for higher accuracies and thus will not be replicating the previously mentioned methodologies, however from the aforementioned papers we conclude that a large dimensionality reduction step will be necessary to efficiently process a large visual data set.

Dimensionality reduction and efficient matching of large amounts of visual data is important within content based image retrieval, Zheng et al. give a very detailed summary of all the significant contributions from past years [10]. Their summary covers three periods in image retrieval: Early methods, SIFT-based methods and Convolutional Neural Network (CNN) based methods. Smeulders et al. present all the contributions of the early methods [11], these methods focused on looking at the color, texture and local geometry of images for retrieval [12, 13]. Not much later the Bag-of-words (BoW) model was proposed as a new method for image retrieval [14]. The advantage of this is that inverted indexes can be used for immediate retrieval of similar images. The BoW model paired with SIFT-descriptors [15] as the feature vectors was used in image retrieval research for years [16, 17, 18, 19, 20]. Since 2012 when the convolutional neural network was introduced [21] research switched to CNN-based methods for image retrieval because they achieved better performance on several image retrieval tasks [22, 23, 24].

3 Image Retrieval

Image retrieval is a subset of the research field Information Retrieval (IR). It investigates the problem where one has a query \mathbf{q} expressing an information need and wants to find the best possible match for this \mathbf{q} in a set of documents \mathbf{D} .

In the case of image retrieval the query consists of an image for which the best matching image should be found in a document set of images. There are two approaches for doing this.

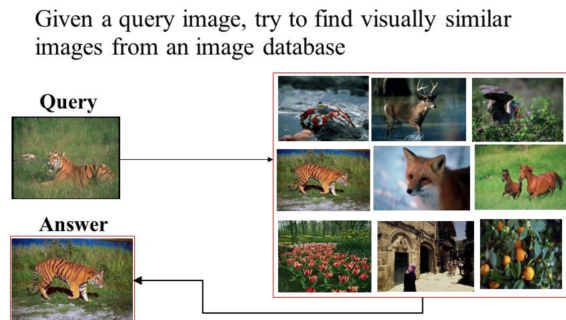


Figure 3: Visual example of image retrieval.

Text-based image retrieval refers to retrieval of images based on textual metadata associated with these images. It matches images based on for instance their titles, keywords and more. The downside of this method is that if there is no metadata available then it needs to be added manually. With the growing sizes of image databases this can become quite inefficient, hence the more focus on content-based image retrieval in past years [25].

Content-based image retrieval is retrieval based on the content of the image rather than the metadata as is the case with concept-based image retrieval. Computer vision is used to evaluate the image similarity, this can be done by looking at colors, shapes, textures and more. The advantage of content based image retrieval is that it does not rely on the metadata of images and thus does not require manual labeling.

3.1 Feature vectors

A feature vector is a vector containing a numeric representation of multiple characteristics of an object. In this case the objects will be images, the frames of a video. To illustrate feature vectors for images, a small example will be given. Consider the following three images of 10 by 10 pixels consisting only of black or white pixels in Figure 4.

If we would like to measure the image similarity between these pictures then we could do this on a pixel-per-pixel basis, this could work for low dimensional images. But comparing HD images consisting of 1920 pixels by 1080 pixels would require massive computation and will require extensive logic to handle small deviations between images. This can be resolved by creating feature vectors to represent the images. A representation for these example images could be the number of black and white pixels in the image. This results in the feature vector $\mathbf{x} = [\# \text{ white pixels}, \# \text{ black pixels}]^T$. This vector can now be represented in 2D-space, and compared to other vectors in 2D space by calculating the distance between other points in that space. The distance is computed by calculating the euclidean distance between each vector.

$$d(\mathbf{p}, \mathbf{q}) = \sqrt{(p_1 - q_1)^2 + \dots + (p_n - q_n)^2}$$



Figure 4: Three images of 10x10 pixels, consisting of only black and white pixels. The resulting feature vectors are $[64, 36]^T$, $[0, 100]^T$ and $[25, 75]^T$ respectively.

In Figure 5 the plotted feature vectors and the distances between them can be viewed. Example images 2 and 3 are the most alike because the distance between them is the smallest.



Figure 5: Example of distance in 2D between feature vectors that represent the example images. According to the chosen features, image 2 and 3 are the most alike.

The similarity of images can be retrieved by computing the distance between their feature vectors. A low distance indicates a higher similarity. Counting the number of black and white pixels for the construction of feature vectors is not a very robust method. Many research in image retrieval has been done on the construction of these feature vectors for achieving the best retrieval results. Part of this research is determining the best method for the construction of the image feature vectors to apply for our problem. We will expand on the methods that are going to be explored.

As mentioned in Section 2, Zheng et al. [10] outline three significant periods in image retrieval, the early methods focused on global descriptors mostly based on color and texture. These global descriptors were not good at handling image changes in illumination, translation, occlusion and truncation. This gave rise to local descriptors based SIFT methods until recently CNN-descriptors became the most popular.

We want to determine whether global or local descriptors are needed. Therefore we take two different global descriptor methods and one one recent local descriptor

method with one of the best performances. The methods are color histograms as a global descriptor, color texture moments [26] as a method that combines the color and texture of an image and lastly CNN-descriptors were implemented as local descriptors.

This is not an exhaustive list of methods, but because computing feature vectors for a large amount of frames is resource intensive, we limited ourselves to these three. If one of the methods achieves promising results then future research can expand on variations of the method or if none of the methods achieve acceptable results, then future research could try other distinct methods.

3.2 Color Histograms

A color histogram is a representation of the distribution of colors in an image. Color histograms are a flexible and low dimensional way of representing images. A color histogram can be computed by counting the number of pixels in a certain color range, the size of this range is variable, called the bin size. A larger bin size results in lower dimensions of the resulting histogram. For example, if one chooses a bin size that is half of the intensity range. The resulting vector would have 6 dimensions, two bins for each color channel. See Figure 6 for an example with such a bin size.




Figure 6: Color histogram with a large binsize (binsize=128) for illustration.

Color histograms are global descriptors because they describe the image on a global level.

3.3 Color Texture Moments

Color texture moments as proposed by Yu et al. [26] are a low-level feature descriptor that integrates both color and texture characteristics of an image. An image is converted from RGB to HSV color space, and on each color map a 2d convolution is applied for each 3x3 filter shown in Figure 7.

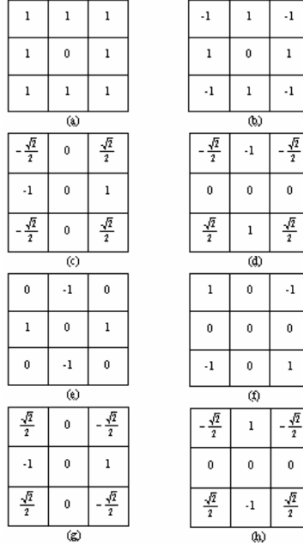


Figure 7: Filters used for computing the color texture channels.

Which results in 8 channels for each color map, 24 channels in total. For each of the resulting channels the first and second color moments, which are the mean and standard deviation, are taken. This results in a 48-dimensional feature vector.

3.4 CNN Features

Convolutional neural networks are very good at image classification and segmentation tasks. However it was found that using the resulting channels after the convolution layers as feature vectors also perform well at image retrieval tasks. There are many ways to construct CNN feature vectors that perform well as described in [10].

We will use the Regional Maximum Activation of Convolutions (R-MAC) method introduced by Tolias et al. [24] as our implementation for the feature vectors. The frames are fed through a pre-trained VGG16 neural network. The last fully connected layers of the network are discarded, and the resulting channels with activations from the convolutional layers are divided up into regions. For each of these regions the Maximum Activations of Convolutions (MAC) feature vector is calculated, and all these feature vectors are then post-processed by applying L2-normalization, PCA-whitening and L2-normalization. Then all the feature vectors are summed into one resulting feature vector and L2-normalization is applied one last time as the final step, resulting in a 512-dimensional vector. For a more extensive description, one can refer to [24].

4 Methodology

We want to explore whether we can detect recurring content unsupervised and if so, which method of feature vector construction scores the best in terms of precision and recall. Initially a non-generalized approach was tried, where a unique method for each content class was used (For example: OCR for textual closing credits, similarity matrices for opening credits and feature based classification for recaps and previews). However, this was quickly found to be ineffective because of a high variety of edge cases in creative non-structured video content.

The characteristic of all the segment classes is that they reoccur either in previous or future episodes. To match frames from one episode with other episodes, feature vectors were constructed from the frames and the distances between these vectors were computed and saved. If a part of the video consecutively has matching feature vectors from previous or future episodes, then that part is labeled as a recurring segment. It is therefore important that the resulting feature vectors do not mismatch between actual recurring content and content that looks alike. Methods chosen for construction of the feature vectors should therefore not result in feature vectors of very high dimensionality, because this has proven to reduce the distance between the farthest and closest points [27], likely diminishing the distinctive property of the feature vectors that is needed.

We will do separate experiments for the bumper detection as they require a different approach for detection because they only contain recurring content from within the episode itself.

All of the experiments were set up in Python and executed on AWS m4.2xlarge instances.

4.1 Data

The data set contains 80 video files originating from 16 different seasons of tv-shows, amounting to around 50 hours of content. Only for three of these seasons do we use the full array of episodes, from the rest of the seasons, only the first three episodes will be looked at to have a more varied data set. Originally these files were in .mxf full broadcast format, this meant each file being 25GB on average. All the files were converted to MP4 files having a width of 320 pixels width and the length to maintain the aspect ratio. The resizing step was done to drastically decrease the computing time needed. The conversion was done with FFmpeg [28].

For each file the start and end timestamps for the recaps, opening credits, closing credits and previews were annotated in HH:MM:SS format in a CSV file, so that it could be loaded and compared effectively.

Table 1 shows the distribution of the different recurring content classes. It can be noted that the classes are not in balance, the recurring content is weighted towards recaps and opening credits. Recurring content consists of 4.9% of the total content, thus with a perfect method, that is the percentage of time we save of people watching the content back-to-back.

	Total seconds	% of recurring	% of total
Recaps	3275	34.8%	1.7%
Opening credits	4669	49.5%	2.5%
Closing credits	690	7.3%	0.4%
Previews	592	6.3%	0.3%
Recurring	9423	-	4.9%
Total	188981	-	-

Table 1: Details of the data that was used to perform the experiments with.

4.2 Feature Vector Construction

Video usually plays at 25 frames per second, frames very close to each other should only have very slight variations. That is why frames every 3 frames or just shot-boundary frames will be taken into account to reduce the computing complexity. We take all the video files of a season $E_x \in T$ and convert each file to a set of feature vectors $S_x = \{v_1, v_2, \dots, v_l\}$ with the function f .

$$f(E_x) = S_x = \{v_1, v_2, \dots, v_l\}$$

The rest of the subsections will expand on different implementations of the function f , how the different types of feature vectors are constructed. We use three different methods for the feature vector construction. These methods were chosen because of their efficient computation and accuracy in image retrieval tasks and to see if there is a difference between global and local descriptors applied to this problem. The methods chosen were color histograms, color texture moments and CNN features.

4.2.1 Color Histograms

For the construction of the color histograms the pixels intensities in each channel (Red, Green, Blue (RGB)) are counted and binned according to a specified bin size. All these bins will be concatenated to result in a feature vector, the size of the feature vector is a result of the chosen bin size.

60 bins are used to represent a single color channel. These are then concatenated into a 1D vector, resulting in a feature vector with a dimensionality of 180. This vector is then L1 normalized, such that:

$$\mathbf{v} = \frac{v_i}{\sum_{i=1}^n |v_i|}$$

A histogram is computed for every 3 frames or shotboundary frame in the video and then stored and saved in a list that preserves the order of the video. This method will be named CH in the results section.

4.2.2 Color Texture Moments

This section will be called CTM in the results section. TODO write completely

4.2.3 CNN Features

For the computation of the CNN feature vectors we use a Github re-implementation written in Python [29]. This section will be called CNN in the results section. TODO write completely

4.3 Matching and Detection

All the resulting lists of feature vectors will be added to the index of an empty Faiss instance. Faiss is a library for efficient similarity search of dense vectors [30, 31] and thus perfectly suited for our task. This library will efficiently handle building the inverted index and nearest neighbor matching. The science that is behind large scale matching of dense vectors is a research area on its own, but outside the scope of this thesis. We use Faiss because it is released as an open source library and therefore easy to use in an implementation.

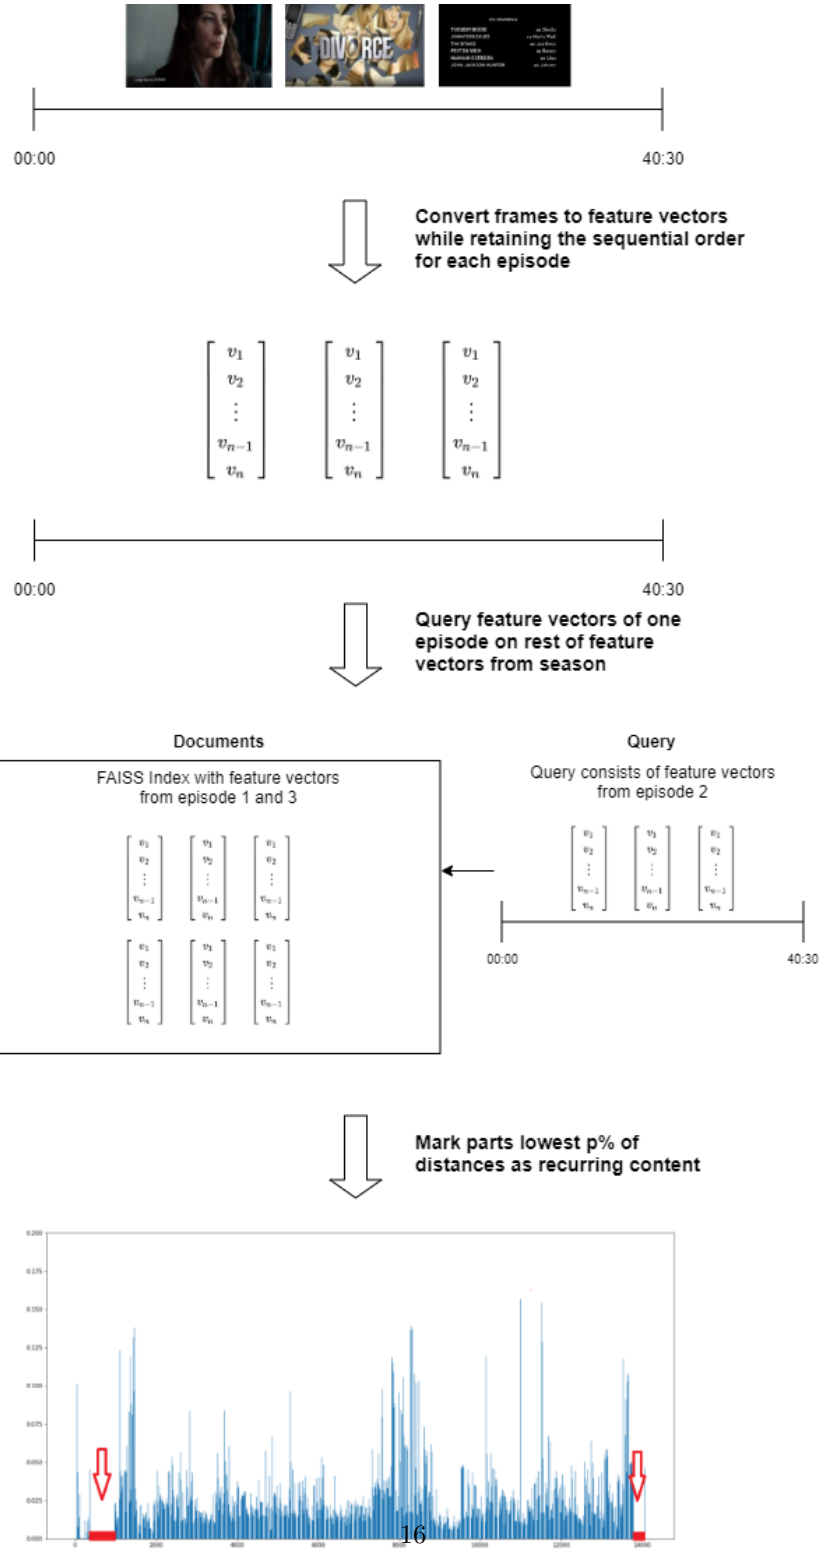


Figure 8: Diagram of the described methodology. The recurring content detection process for episode 2 is portrayed.

4.3.1 Recurring Content

To mark a part of the video as recurring content, we do the following steps: The frames of interest are converted to feature vectors while retaining the sequential order of episodes. Then the episode is queried on the feature vectors of the rest of the episodes, this results in a vector with distances for the best matches. We mark the locations of the indexes where these distances are lower than the percentile $p \in \{5\%, 7.5\%, 10\%\}$ of the lowest-distances vector. These are our initial detections. Because these detections can still contain noise, we do the following post-processing steps: Detections are merged into one if they are less than 10 seconds apart and from the detections we only keep the two largest within the first 20% of the video or we keep the detection if it ends in the last 15 seconds of the video. The rationale behind this is that all recurring content that we want to detect is either in the first or last part of the video. See Figure 8 for a visual depiction of the described methodology.

4.3.2 Bumpers

The methodology for detecting the bumpers within video content is different. Because a bumper only contains recurring content from within the video itself. It shows shots that will be shown after the break, these shots may not reoccur in other episodes within a season. Therefore we cannot use the same methodology as described before for detection.

For the detection of bumpers we take a sliding window of 20 seconds of video frames, this will be the query. The index will consist of the remaining frames in the video, for every frame of interest the resulting distances of the two nearest neighbors will be saved and then this process is repeated for the next 20 seconds in the video all the way to the end.

After all the closest in-episode nearest neighbors are computed for each frame then from these results bumpers can be detected. We found that a global threshold for determining matches does not work well, probably because of the fact that bumpers are very short. Therefore we will use the ratio test described by Lowe [15] for deciding on whether there is a match. With the ratio test the distance between the closest neighbor and second-closest neighbor is taken into account. The ratio of the distance of the closest neighbor to the distance of the second-closest neighbor determines whether there is a correct match, if the ratio is low then there is a higher chance for the match to be correct. In [15] they classify a match to be correct if the ratio of the distances is equal to 0.8 or lower, we experimented with a ratio $r \in \{0.5, 0.6, 0.7, 0.8\}$ to find the optimal ratio for this problem.

5 Results

This section contains the results for the experiments of detecting the recurring content (recaps, opening/closing credits and previews) and the experiments of the bumper detection.

5.1 Recurring Content

The results were evaluated as a retrieval problem. By doing so, the measures precision and recall could be used to describe several aspects of each method. Each second of recurring content is marked as a relevant instance and if recurring content gets detected, then those are relevant detected seconds, by doing so we can calculate precision and recall as following:

$$\text{precision} = \frac{\text{detected relevant seconds}}{\text{total detected seconds}}$$

$$\text{recall} = \frac{\text{detected relevant seconds}}{\text{total relevant seconds}}$$

Because it is possible for the manual annotation to also be slightly inaccurate, we mark detections that are within a 2 second margin of the relevant section, to also be relevant. Figure 9 depicts how the aforementioned precision and recall formulas apply on the detections and ground truth of a single episode. For the calculation of the total results on all episodes, the 'seconds' variables of all episodes are summed together and then precision and recall are calculated. All the results can be viewed in Table 2. Precision-recall curves were also plotted for each recurring content class separately, these can be viewed in Figure 10.




Figure 9: Visual example of the described evaluation methodology.

Method	P	R	P	R	P	R
Uniform Sampling						
	5% lowest		7.5% lowest		10% lowest	
CH	0.884	0.687	0.830	0.812	0.786	0.855
CTM	0.883	0.681	0.841	0.796	0.796	0.853
CNN	0.854	0.733	0.808	0.838	0.748	0.892
Shotboundaries						
CH	0.904	0.529	0.852	0.638	0.763	0.707
CTM	0.856	0.515	0.830	0.652	0.771	0.724
CNN	0.837	0.522	0.800	0.635	0.725	0.713

Table 2: Results of the recurrent content detection at varying values for the lowest percentiles of distances.

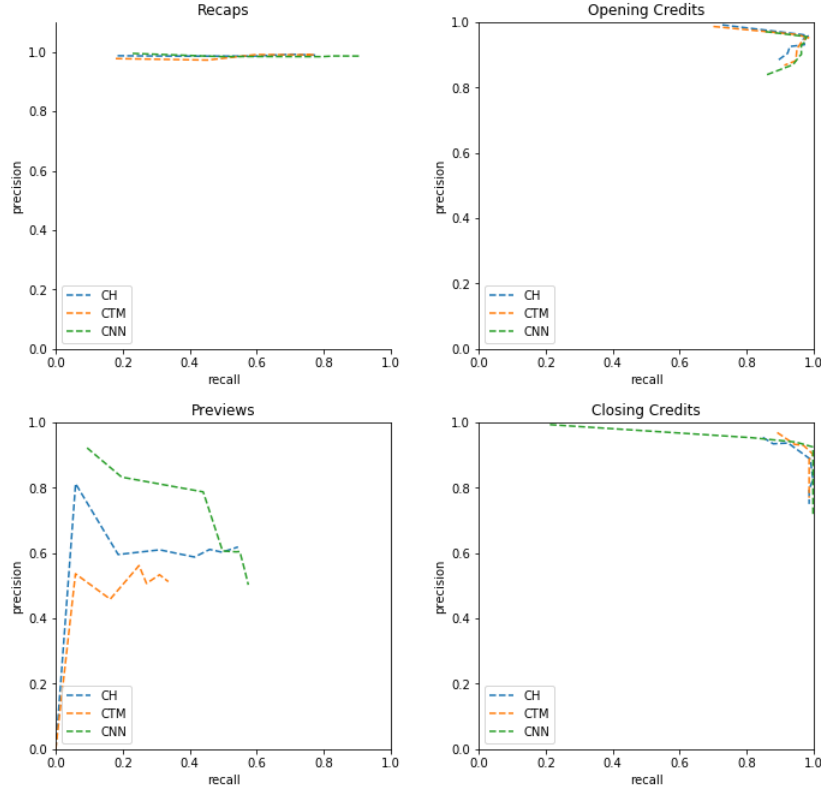


Figure 10: Precision-Recall curves for each recurrent content class for the percentile values in [2.5, 20]. Data points that are in the upper right indicate good performance.

5.2 Bumpers

Because bumpers are short segments, the bumpers were evaluated in a different manner. We do not calculate precision and recall based on numbers of retrieved relevant seconds but whether a detection has overlap with a bumper. If a detection has overlap with a bumper, then this detection is marked as relevant. Precision and recall will thus be calculated according to the following formulas:

$$\text{precision} = \frac{\text{detections with overlap}}{\text{total detections}}$$

$$\text{recall} = \frac{\text{detections with overlap}}{\text{total bumpers in episode}}$$

Figure 11 depicts how the precision and recall is calculated for the bumper detection. All the results of these experiments can be viewed in Table 3

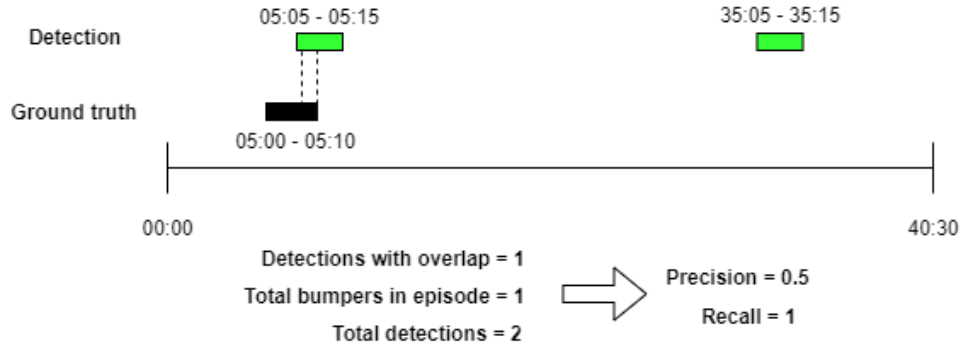


Figure 11: Visual example of the described evaluation methodology for bumper detection.

Method	P	R	P	R	P	R	P	R
	0.5 ratio		0.6 ratio		0.7 ratio		0.8 ratio	
CH180	0.195	0.312	0.203	0.562	0.129	0.646	0.085	0.792
CTM	0.145	0.542	0.089	0.667	0.072	0.833	0.072	0.958
CNN	0.611	0.917	0.484	0.938	0.288	0.979	0.081	0.979

Table 3: Results of the experiments on the bumper detection for varying values of the ratio test.

6 Discussion

The results show that for the recurrent segment detection, the results for each different method for feature vector construction are quite close. Meaning that the way the

feature vectors are constructed does not improve the results significantly when looking at the total results. However, when looking at the scores for each class separately in Figure 10, it can be noticed that the CNN-method does better on the recaps and previews retrieval tasks. Possible explanations for this could be that the CNN-method is a local descriptor as opposed to a global one, therefore being able to handle the more tough cases of recaps and previews. Another explanation could be that higher dimensionality vectors are more suited, because the CNN vectors had the highest dimensionality.

Table 2 presents that using the shotboundaries results in lower recall scores compared to the uniform sampling of taking a frame every 3 frames. An possible explanation is that the shotboundary frames sometimes mismatch because the content shown does not consist of full shots.

Moreover, it is noticed that in total, the detection results are decent but not extraordinary, meaning that there is room for improvement in the methodology.

From the results of the bumper detection we notice there is a difference in results for each method. The CNN-based method outperforms the other methods significantly. An explanation for this could be that again the CNN feature vectors are local descriptors, it could be possible that the ratio test favors these kind of descriptors as it was first proposed in a paper about a local descriptor.

With these results we set the first steps towards a fully automated solution for detection of recurrent content. Many research has been done on broadcast streams, this thesis presents the first results on VOD content. Based on the results of the experiments it becomes clear that feature vectors constructed via a CNN have the most potential for the application to our problem.

The experiments were mostly performed on the first three episodes of a single season. However, an end-solution would take a whole season into account. Also, sometimes the recurring content spans across different seasons (recapping the previous season in the first episode for example), our experiments did not take this attribute into account.

The bumper detection is based on how RTL inserts its bumpers into content. With our experiments, we did not test if our methodology also applies to the bumpers of other publishers of commercial content.

Ideally, future studies could be performed on a larger dataset consisting of full seasons of several shows to further prove the validity of this method and these future studies could try variations of the methodology and different ways of constructing the CNN feature vectors. Also the experiment could be repeated with data where the recurrent content classes are more in balance to verify the results.

7 Conclusion

In this thesis we explored how well an image retrieval based approach could detect recurring content in a TV-show in an unsupervised manner. Experiments were performed using three different kind of feature vectors. Based upon these experiments we can conclude that an image retrieval based approach definitely has potential for the

unsupervised detection of recurrent content in TV-Shows. The experiments showed that the three different methods all get reasonable results. Overall, it is found that the CNN-based method has the most potential. Using this method for a metadata labeling task will still require manual supervision. Future research is needed to improve upon the methodology or to other implementations of CNN-based feature vectors. This thesis presents novel research for the unexplored research area of recurring content detection that is focused on VOD content.

8 References

- [1] R. W. Lienhart, “Comparison of automatic shot boundary detection algorithms,” in *Storage and Retrieval for Image and Video Databases VII*, vol. 3656, pp. 290–302, International Society for Optics and Photonics, 1998.
- [2] H. Shao, Y. Qu, and W. Cui, “Shot boundary detection algorithm based on hsv histogram and hog feature,” in *5th International Conference on Advanced Engineering Materials and Technology*, pp. 951–957, 2015.
- [3] R. Lienhart, C. Kuhmünch, and W. Effelsberg, “On the detection and recognition of television commercials,” 1997.
- [4] J. M. Gauch and A. Shivadas, “Finding and identifying unknown commercials using repeated video sequence detection,” *Computer Vision and Image Understanding*, vol. 103, no. 1, pp. 80–88, 2006.
- [5] M. Covell, S. Baluja, and M. Fink, “Advertisement detection and replacement using acoustic and visual repetition,” in *Multimedia Signal Processing, 2006 IEEE 8th workshop on*, pp. 461–466, IEEE, 2006.
- [6] J. Wang, L. Duan, Q. Liu, H. Lu, and J. S. Jin, “A multimodal scheme for program segmentation and representation in broadcast video streams,” *IEEE Transactions on Multimedia*, vol. 10, no. 3, pp. 393–408, 2008.
- [7] C. Herley, “Argos: Automatically extracting repeating objects from multimedia streams,” *IEEE Transactions on multimedia*, vol. 8, no. 1, pp. 115–129, 2006.
- [8] Y. Benezeth and S.-A. Berrani, “Unsupervised credit detection in tv broadcast streams,” in *Multimedia (ISM), 2010 IEEE International Symposium on*, pp. 175–182, IEEE, 2010.
- [9] A. E. Abduraman, S.-A. Berrani, and B. Merialdo, “An unsupervised approach for recurrent tv program structuring,” in *Proceedings of the 9th European Conference on Interactive TV and Video*, pp. 123–126, ACM, 2011.
- [10] L. Zheng, Y. Yang, and Q. Tian, “SIFT meets CNN: A decade survey of instance retrieval,” *IEEE transactions on pattern analysis and machine intelligence*, vol. 40, no. 5, pp. 1224–1244, 2018.

- [11] A. W. Smeulders, M. Worring, S. Santini, A. Gupta, and R. Jain, “Content-based image retrieval at the end of the early years,” *IEEE Transactions on Pattern Analysis & Machine Intelligence*, no. 12, pp. 1349–1380, 2000.
- [12] H. Yu, M. Li, H.-J. Zhang, and J. Feng, “Color texture moments for content-based image retrieval,” in *Proceedings. International Conference on Image Processing*, vol. 3, pp. 929–932, IEEE, 2002.
- [13] B. S. Manjunath and W.-Y. Ma, “Texture features for browsing and retrieval of image data,” *IEEE Transactions on pattern analysis and machine intelligence*, vol. 18, no. 8, pp. 837–842, 1996.
- [14] J. Sivic and A. Zisserman, “Video google: A text retrieval approach to object matching in videos,” in *null*, p. 1470, IEEE, 2003.
- [15] D. G. Lowe, “Distinctive image features from scale-invariant keypoints,” *International journal of computer vision*, vol. 60, no. 2, pp. 91–110, 2004.
- [16] D. Nister and H. Stewenius, “Scalable recognition with a vocabulary tree,” in *2006 IEEE Computer Society Conference on Computer Vision and Pattern Recognition (CVPR’06)*, vol. 2, pp. 2161–2168, Ieee, 2006.
- [17] J. Philbin, O. Chum, M. Isard, J. Sivic, and A. Zisserman, “Object retrieval with large vocabularies and fast spatial matching,” in *2007 IEEE conference on computer vision and pattern recognition*, pp. 1–8, IEEE, 2007.
- [18] H. Jegou, M. Douze, and C. Schmid, “Hamming embedding and weak geometric consistency for large scale image search,” in *European conference on computer vision*, pp. 304–317, Springer, 2008.
- [19] H. Jégou, M. Douze, C. Schmid, and P. Pérez, “Aggregating local descriptors into a compact image representation,” in *CVPR 2010-23rd IEEE Conference on Computer Vision & Pattern Recognition*, pp. 3304–3311, IEEE Computer Society, 2010.
- [20] H. Jegou, F. Perronnin, M. Douze, J. Sánchez, P. Perez, and C. Schmid, “Aggregating local image descriptors into compact codes,” *IEEE transactions on pattern analysis and machine intelligence*, vol. 34, no. 9, pp. 1704–1716, 2012.
- [21] A. Krizhevsky, I. Sutskever, and G. E. Hinton, “Imagenet classification with deep convolutional neural networks,” in *Advances in neural information processing systems*, pp. 1097–1105, 2012.
- [22] A. Babenko, A. Slesarev, A. Chigorin, and V. Lempitsky, “Neural codes for image retrieval,” in *European conference on computer vision*, pp. 584–599, Springer, 2014.
- [23] J. Yue-Hei Ng, F. Yang, and L. S. Davis, “Exploiting local features from deep networks for image retrieval,” in *Proceedings of the IEEE conference on computer vision and pattern recognition workshops*, pp. 53–61, 2015.

- [24] G. Tolias, R. Sivic, and H. Jégou, “Particular object retrieval with integral max-pooling of cnn activations,” *arXiv preprint arXiv:1511.05879*, 2015.
- [25] F. Rajam and S. Valli, “A survey on content based image retrieval,” *Life Science Journal*, vol. 10, no. 2, pp. 2475–2487, 2013.
- [26] H. Yu, M. Li, H.-J. Zhang, and J. Feng, “Color texture moments for content-based image retrieval,” in *Proceedings. International Conference on Image Processing*, vol. 3, pp. 929–932, IEEE, 2002.
- [27] K. Beyer, J. Goldstein, R. Ramakrishnan, and U. Shaft, “When is “nearest neighbor” meaningful?,” in *International conference on database theory*, pp. 217–235, Springer, 1999.
- [28] “FFmpeg.” <http://ffmpeg.org>.
- [29] “Re-implementation of regional maximum activations of convolutions (RMAC) feature extractor for Keras.” https://github.com/noagarcia/keras_rmac.
- [30] J. Johnson, M. Douze, and H. Jégou, “Billion-scale similarity search with GPUs,” *arXiv preprint arXiv:1702.08734*, 2017.
- [31] “Faiss, a library for efficient similarity search and clustering of dense vectors - Facebook Research.” <https://github.com/facebookresearch/faiss>.